Numerical Optimization

Lecture Notes #9 — Trust-Region Methods
Global Convergence and Enhancements

Fall 2024

Outline

Recap: — Iterative "Nearly Exact" Solution of the Subproblem

Last time we looked at **nearly exact solution** of the subproblem

$$\min_{\bar{\mathbf{p}} \in T_k} m_k(\bar{\mathbf{p}}) = \min_{\bar{\mathbf{p}} \in T_k} f(\bar{\mathbf{x}}_k) + \bar{\mathbf{p}}^T \nabla f(\bar{\mathbf{x}}_k) + \frac{1}{2} \bar{\mathbf{p}}^T B_k \bar{\mathbf{p}}$$

This approach is viable for problems with few degrees of freedom, *e.g.*
$T_k \subseteq \mathbb{R}^n$, n "small." Where "small" means that the **unitary
diagonalization** $Q_k \Lambda_k Q_k^T = B_k$ is computable in a "reasonable" amount
of time.

From a theoretical characterization of the exact problem, we derived an
algorithm which finds a nearly exact solution at a cost per iteration
approximately **three** times that of dogleg and 2D-subspace minimization.

The scheme was based on a 1-D Newton iteration (with some clever
tricks), and some careful analysis of special (hard) cases.

On Today's Menu

We wrap up the first pass of Trust Region methods —

– We briefly discuss global convergence properties for trust region methods.

– We look at some theorems, but leave the proofs as "exercises."

– For second order $(B_k \neq \nabla^2 f(\bar{\mathbf{x}}_k))$ models we can show convergence to a stationary point.

– For trust-region Newton methods $(B_k = \nabla^2 f(\bar{\mathbf{x}}_k))$ models we can show convergence to a point where the second order necessary conditions hold.

– We look at modifications for poorly scaled problems, as well as the use of non-spherical trust regions.

Theorem (Second Order Necessary Conditions)

If $\bar{\mathbf{x}}^*$ is a local minimizer of f and $\nabla^2 f$ is continuous in an open neighborhood of $\bar{\mathbf{x}}^*$, then $\nabla f(\bar{\mathbf{x}}^*) = 0$ and $\nabla^2 f(\bar{\mathbf{x}}^*)$ is positive semi-definite.

Global Convergence
Global Convergence...
Enhancements

Tool #1 — A Lemma: The Cauchy Point
Tool #2 — A Theorem
Recall: The Trust Region Algorithm

Global Convergence: Tool #1 — A Lemma

Recall: The trust-region subproblem is

$$\bar{\mathbf{p}}_k = \underset{\|\bar{\mathbf{p}}\| \leq \Delta_k}{\arg \min} \, m_k(\bar{\mathbf{p}}) = \underset{\|\bar{\mathbf{p}}\| \leq \Delta_k}{\arg \min} \, f(\bar{\mathbf{x}}_k) + \bar{\mathbf{p}}^T \nabla f(\bar{\mathbf{x}}_k) + \frac{1}{2}\bar{\mathbf{p}}^T B_k \bar{\mathbf{p}}.$$

The following lemma gives us a lower bound for the decrease in the model at the Cauchy point:

Lemma (Cauchy point descent)

The Cauchy point $\bar{\mathbf{p}}_k^c$ satisfies

$$m_k(\bar{\mathbf{0}}) - m_k(\bar{\mathbf{p}}_k^c) \geq \frac{1}{2}\|\nabla f(\bar{\mathbf{x}}_k)\| \, \min\left[\Delta_k, \frac{\|\nabla f(\bar{\mathbf{x}}_k)\|}{\|B_k\|}\right].$$

Global Convergence
Global Convergence...
Enhancements

Tool #1 — A Lemma: The Cauchy Point
Tool #2 — A Theorem
Recall: The Trust Region Algorithm

Proof of Lemma   The Cauchy Point

We recall the explicit expressions for the Cauchy point (from lecture 7)

$$
\begin{cases}
\quad \bar{\mathbf{p}}_k^c \;=\; -\tau_k \dfrac{\Delta_k}{\|\nabla f(\bar{\mathbf{x}}_k)\|}\nabla f(\bar{\mathbf{x}}_k) \\[2mm]
\text{where} \\[2mm]
\quad \tau_k \;=\;
\begin{cases}
1 & \text{if } \nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k) \le 0 \\[2mm]
\min\left(1, \dfrac{\|\nabla f(\bar{\mathbf{x}}_k)\|^3}{\Delta_k \nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k)}\right) & \text{otherwise}
\end{cases}
\end{cases}
$$

Figure: The three possible scenarios for selection of τ.

Global Convergence
Global Convergence...
Enhancements

Tool #1 — A Lemma: The Cauchy Point
Tool #2 — A Theorem
Recall: The Trust Region Algorithm

Proof of Lemma   Case#1

Case#1 $(\nabla f(\bar{\mathbf{x}}_k) B_k \nabla f(\bar{\mathbf{x}}) \leq 0)$:

In this scenario $m_k(\bar{\mathbf{p}}_k^c) - m_k(\bar{\mathbf{0}}) =$

$$
\begin{aligned}
&= m_k\left(-\Delta_k \frac{\nabla f(\bar{\mathbf{x}}_k)}{\|\nabla f(\bar{\mathbf{x}}_k)\|}\right) - m_k(\bar{\mathbf{0}}) \\
&= -\Delta_k \|\nabla f(\bar{\mathbf{x}}_k)\| + \frac{1}{2}\frac{\Delta_k^2}{\|\nabla f(\bar{\mathbf{x}}_k)\|^2} \underbrace{\nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k)}_{\leq 0} \\
&\leq -\Delta_k \|\nabla f(\bar{\mathbf{x}}_k)\| \\
&\leq -\|\nabla f(\bar{\mathbf{x}}_k)\| \min\left(\Delta_k, \frac{\|\nabla f(\bar{\mathbf{x}}_k)\|}{\|B_k\|}\right)
\end{aligned}
$$

Hence,

$$
m_k(\bar{\mathbf{0}}) - m_k(\bar{\mathbf{p}}_k^c) \geq \|\nabla f(\bar{\mathbf{x}}_k)\| \min\left(\Delta_k, \frac{\|\nabla f(\bar{\mathbf{x}}_k)\|}{\|B_k\|}\right) \geq \frac{1}{2}\|\nabla f(\bar{\mathbf{x}}_k)\| \min\left(\Delta_k, \frac{\|\nabla f(\bar{\mathbf{x}}_k)\|}{\|B_k\|}\right)
$$

## Proof of Lemma

Case#2

Case#2 $(\nabla f(\bar{\mathbf{x}}_k)B_k\nabla f(\bar{\mathbf{x}}) > 0$, and $\frac{\|\nabla f(\bar{\mathbf{x}}_k)\|^3}{\Delta_k \nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k)} \leq 1)$:

In this scenario the Cauchy point is in the interior of the trust region, and $m_k(\bar{\mathbf{p}}_k^c) - m_k(\bar{\mathbf{0}}) =$

$$
\begin{aligned}
&= -\frac{\|\nabla f(\bar{\mathbf{x}}_k)\|^4}{\nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k)} + \frac{1}{2}\frac{\|\nabla f(\bar{\mathbf{x}}_k)\|^4}{(\nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k))^2}\nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k) \\
&= -\frac{1}{2}\frac{\|\nabla f(\bar{\mathbf{x}}_k)\|^4}{\nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k)} \\
&\leq -\frac{1}{2}\frac{\|\nabla f(\bar{\mathbf{x}}_k)\|^4}{\|B_k\|\,\|\nabla f(\bar{\mathbf{x}}_k)\|^2} = -\frac{1}{2}\frac{\|\nabla f(\bar{\mathbf{x}}_k)\|^2}{\|B_k\|} \\
&\leq -\frac{1}{2}\|\nabla f(\bar{\mathbf{x}}_k)\|\min\left(\Delta_k, \frac{\|\nabla f(\bar{\mathbf{x}}_k)\|}{\|B_k\|}\right)
\end{aligned}
$$

Use the minus sign to flip the inequality, and we're there!

Proof of Lemma

Case#3

Case#3 ($\nabla f(\bar{\mathbf{x}}_k) B_k \nabla f(\bar{\mathbf{x}}) > 0$, and $\frac{\|\nabla f(\bar{\mathbf{x}}_k)\|^3}{\Delta_k \nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k)} > 1$):

We note that in this scenario $\nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k) < \frac{\|\nabla f(\bar{\mathbf{x}}_k)\|^3}{\Delta_k}$, and $m_k(\bar{\mathbf{p}}_k^c) - m_k(\bar{\mathbf{0}}) =$

$$
\begin{aligned}
&= -\frac{\Delta_k}{\|\nabla f(\bar{\mathbf{x}}_k)\|}\|\nabla f(\bar{\mathbf{x}}_k)\|^2 + \frac{1}{2}\frac{\Delta_k^2}{\|\nabla f(\bar{\mathbf{x}}_k)\|^2}\nabla f(\bar{\mathbf{x}}_k)^T B_k \nabla f(\bar{\mathbf{x}}_k) \\
&\leq -\Delta_k\|\nabla f(\bar{\mathbf{x}}_k)\| + \frac{1}{2}\frac{\Delta_k^2}{\|\nabla f(\bar{\mathbf{x}}_k)\|^2}\frac{\|\nabla f(\bar{\mathbf{x}}_k)\|^3}{\Delta_k} \\
&= -\frac{1}{2}\Delta_k\|\nabla f(\bar{\mathbf{x}}_k)\| \\
&\leq -\frac{1}{2}\|\nabla f(\bar{\mathbf{x}}_k)\| \min\left(\Delta_k, \frac{\|\nabla f(\bar{\mathbf{x}}_k)\|}{\|B_k\|}\right)
\end{aligned}
$$

Use the minus sign to flip the inequality, and we're there!

Global Convergence
Global Convergence...
Enhancements

Tool #1 — A Lemma: The Cauchy Point
Tool #2 — A Theorem
Recall: The Trust Region Algorithm

Global Convergence: Tool #2 — A Theorem

Theorem

Let $\bar{\mathbf{p}}_k$ be any vector, $\|\bar{\mathbf{p}}_k\| \leq \Delta_k$, such that

$$m_k(\bar{\mathbf{0}}) - m_k(\bar{\mathbf{p}}_k) \geq c_2(m_k(\bar{\mathbf{0}}) - m_k(\bar{\mathbf{p}}_k^c))$$

then

$$m_k(\bar{\mathbf{0}}) - m_k(\bar{\mathbf{p}}_k) \geq \frac{c_2}{2}\|\nabla f(\bar{\mathbf{x}}_k)\| \min\left[\Delta_k, \frac{\|\nabla f(\bar{\mathbf{x}}_k)\|}{\|B_k\|}\right].$$

Both the dogleg, and 2-D subspace minimization algorithms (as well as Steihaug's algorithm) fall into this category, with $c_2 = 1$, since they all produce $\bar{\mathbf{p}}_k$ which give at least as much descent as the Cauchy point, *i.e.* $m_k(\bar{\mathbf{p}}_k) \leq m_k(\bar{\mathbf{p}}_k^c)$.

We are going to use this result to show convergence for the trust region algorithm (see next slide).

Global Convergence
Global Convergence...
Enhancements

Tool #1 — A Lemma: The Cauchy Point
Tool #2 — A Theorem
Recall: The Trust Region Algorithm

The Trust Region Algorithm

Algorithm: Trust Region

```
[ 1] Set k = 1, Δ̂ > 0, Δ₀ ∈ (0, Δ̂), and η ∈ [0, ¼]
[ 2] While optimality condition not satisfied
[ 3]    Get p̄ₖ (approximate solution)
[ 4]    Evaluate ρₖ
[ 5]    if ρₖ < ¼
[ 6]       Δₖ₊₁ = ¼Δₖ
[ 7]    else
[ 8]       if ρₖ > ¾ and ‖p̄ₖ‖ = Δₖ
[ 9]          Δₖ₊₁ = min(2Δₖ, Δ̂)
[10]       else
[11]          Δₖ₊₁ = Δₖ
[12]       endif
[13]    endif
[14]    if ρₖ > η
[15]       x̄ₖ₊₁ = x̄ₖ + p̄ₖ
[16]    else
[17]       x̄ₖ₊₁ = x̄ₖ
[18]    endif
[19]    k = k + 1
[20] End-While
```

Convergence to Stationary Points

Case $\eta = 0$

accept any step which produces descent in f — we can show that the sequence of gradients $\{\nabla f(\bar{\mathbf{x}}_k)\}$ has a **limit point** at zero.

Case $\eta > 0$

accept a step only if the decrease in f is at least some fixed fraction of the predicted decrease — we can show the stronger result $\{\nabla f(\bar{\mathbf{x}}_k)\} \to \bar{\mathbf{0}}$.

In order for the proof(s) to work, we must assume that the model Hessians B_k are uniformly bounded, *i.e.* $\|B_k\| \leq \beta$, and that f is bounded below on the levelset $\{\bar{\mathbf{x}} \in \mathbb{R}^n : f(\bar{\mathbf{x}}) \leq f(\bar{\mathbf{x}}_0)\}$.

The trust-region bound can be relaxed so that the results hold as long as the solution to the subproblems satisfy

$$\|\bar{\mathbf{p}}_k\| \leq \gamma \Delta_k, \quad \text{for some constant } \gamma \geq 1.$$

Convergence to Stationary Points: $\eta = 0$

Theorem

Let $\eta = 0$ in the trust region algorithm. Suppose that $\|B_k\| \leq \beta$ for some constant β, that f is continuously differentiable and bounded below on the bounded set $\{\bar{\mathbf{x}} \in \mathbb{R}^n : f(\bar{\mathbf{x}}) \leq f(\bar{\mathbf{x}}_0)\}$, and that all approximate solutions to the trust-region subproblem satisfy the inequalities

$$m_k(\bar{\mathbf{0}}) - m_k(\bar{\mathbf{p}}_k) \geq c_1 \|\nabla f(\bar{\mathbf{x}}_k)\| \, \min\left[\Delta_k, \frac{\|\nabla f(\bar{\mathbf{x}}_k)\|}{\|B_k\|}\right],$$

and

$$\|\bar{\mathbf{p}}_k\| \leq \gamma \Delta_k,$$

for some positive constants c_1 and γ. **Then** *we have*

$$\liminf_{k \to \infty} \|\nabla f(\bar{\mathbf{x}}_k)\| = 0.$$

Convergence to Stationary Points: $\eta > 0$

Theorem

Let $\eta \in \left(0, \frac{1}{4}\right)$ in the trust region algorithm. Suppose that $\|B_k\| \leq \beta$ for some constant β, that f is Lipschitz continuously differentiable and bounded below on the bounded set $\{\bar{\mathbf{x}} \in \mathbb{R}^n : f(\bar{\mathbf{x}}) \leq f(\bar{\mathbf{x}}_0)\}$, and that all approximate solutions to the trust-region subproblem satisfy the inequalities

$$m_k(\bar{\mathbf{0}}) - m_k(\bar{\mathbf{p}}_k) \geq c_1 \|\nabla f(\bar{\mathbf{x}}_k)\| \, \min\left[\Delta_k, \frac{\|\nabla f(\bar{\mathbf{x}}_k)\|}{\|B_k\|}\right].$$

and

$$\|\bar{\mathbf{p}}_k\| \leq \gamma \Delta_k$$

for some positive constants c_1 and γ. **Then** *we have*

$$\lim_{k \to \infty} \nabla f(\bar{\mathbf{x}}_k) = \bar{\mathbf{0}}.$$

Proofs: Convergence to Stationary Points

The complete proofs are in $\texttt{NW}^{\text{1st}}$ pp.90–91, and pp.92–93; or $\texttt{NW}^{\text{2nd}}$ pp.80–82, and pp.82–83.

The proofs are based on manipulation of ρ — the ratio of actual (objective) reduction and predicted (model) reduction; Taylor's theorem; then deriving a contradiction from the supposition $\|\nabla f(\bar{\mathbf{x}}_k)\| \geq \epsilon$ using careful selection of scalings and bounds for Δ_k.

Definition (lim sup and lim inf)

Let $\{s_n\}$ be a sequence of real numbers. Let E be the set of values x so that $s_{n_k} \to x$ for some subsequence $\{s_{n_k}\}$. This set E contains all sub-sequential limits, plus possibly $\pm\infty$; let

$$s^* = \sup E, \quad s_* = \inf E$$

The values s^* and s_* are the upper and lower limits of $\{s_n\}$, and we use the notation

$$\limsup_{n \to \infty} s_n = s^*, \quad \liminf_{n \to \infty} s_n = s_*$$

Convergence: Iterative "Nearly Exact" Solutions $\bar{\mathbf{p}}_k^*$, for Trust-Region Newton

Theorem ($\text{NW}^{2\text{nd}}$ p.92, proof in Moré & Sorensen (1983))

Let $\eta \in \left(0, \frac{1}{4}\right)$ in the algorithm on slide 11, let $B_k = \nabla^2 f(\bar{\mathbf{x}}_k)$, and suppose that $\bar{\mathbf{p}}_k$ at each iteration satisfy

$$m_k(\bar{\mathbf{0}}) - m_k(\bar{\mathbf{p}}_k) \geq c_1(m_k(\bar{\mathbf{0}}) - m_k(\bar{\mathbf{p}}_k^*)),$$

and $\|\bar{\mathbf{p}}_k\| \leq \gamma \Delta_k$, for some positive constant γ, and $c_1 \in (0, 1]$. **Then**

$$\lim_{k \to \infty} \|\nabla f(\bar{\mathbf{x}}_k)\| = 0.$$

If, in addition, the set $\{\bar{\mathbf{x}} \in \mathbb{R}^n : f(\bar{\mathbf{x}}) \leq f(\bar{\mathbf{x}}_0)\}$ is compact, then **either** *the algorithm terminates at a point $\bar{\mathbf{x}}_k$ at which the second order necessary conditions for a local minimum hold,* **or** *$\{\bar{\mathbf{x}}_k\}$ has a limit point $\bar{\mathbf{x}}^* \in \{\bar{\mathbf{x}} \in \mathbb{R}^n : f(\bar{\mathbf{x}}) \leq f(\bar{\mathbf{x}}_0)\}$ at which the conditions hold.*

Enhancement: Scaling — The Problem

As we have seen before (in the context of steepest descent / line-search), **scaling** (ill-conditioning) can cause problems. — If the objective is more sensitive to changes in one variable than other, the contour lines stretch out to be narrow ellipses (in 2D).

Clearly, a circular trust-region may be quite limiting in this scenario. — The radius is limited by the sensitive variable.

Enhancement: Scaling — The Solution

The solution to the problem of poor scaling is to use **elliptical** trust regions. We define a diagonal scaling matrix

$$D = \operatorname{diag}(d_1, d_2, \ldots, d_n), \quad d_i > 0.$$

Then, the constraint $\|D\bar{\mathbf{p}}\| \leq \Delta$ defines an elliptical trust region, and we get the following scaled trust-region subproblem:

$$\min_{\bar{\mathbf{p}} \in \mathbb{R}^n \,:\, \|D\bar{\mathbf{p}}\| \leq \Delta_k} f(\bar{\mathbf{x}}_k) + \bar{\mathbf{p}}^T \nabla f(\bar{\mathbf{x}}_k) + \frac{1}{2}\bar{\mathbf{p}}^T B_k \bar{\mathbf{p}}.$$

The scaling matrix can be built using information about the gradient $\nabla f(\bar{\mathbf{x}}_k)$ and the Hessian $\nabla^2 f(\bar{\mathbf{x}}_k)$ along the solution path. — We can allow $D = D_k$ to change from iteration to iteration.

All our analysis/algorithms still work with scaling added — but we get factors of D^{-2}, D^{-1}, D, and D^2 in our expressions.

## Feature: Non-Euclidean Trust Regions   1 of 4

Figure: Illustration of (unscaled) trust region boundaries for, from left-to-right: $\|\bar{\mathbf{p}}\|_2 \leq \Delta_k$, $\|\bar{\mathbf{p}}\|_1 \leq \Delta_k$, $\|\bar{\mathbf{p}}\|_4 \leq \Delta_k$, and $\|\bar{\mathbf{p}}\|_\infty \leq \Delta_k$.

Most of the time using trust regions based on norms with $q \neq 2$:

$$\|\bar{\mathbf{p}}\|_q \leq \Delta_k \text{ (unscaled)}, \quad \|D\bar{\mathbf{p}}\|_q \leq \Delta_k \text{ (scaled)}$$

cause us a giant head-ache. There are however some situations when such regions come in handy...

## Feature: Non-Euclidean Trust Regions   2 of 4

Figure: Illustration of (unscaled) trust region boundaries for, from left-to-right: $\|\bar{\mathbf{p}}\|_1 \leq \Delta_k$, $\|\bar{\mathbf{p}}\|_{\frac{1}{2}} \leq \Delta_k$, $\|\bar{\mathbf{p}}\|_{\frac{1}{4}} \leq \Delta_k$, and $\|\bar{\mathbf{p}}\|_{\frac{1}{8}} \leq \Delta_k$.

Using $q < 1$ leads to non-convex trust regions, which may be a bit of a pain?!?

This may, however, be useful/necessary for non-convex optimization problems.

For **constrained** problems, *e.g.*

$$\min_{\bar{\mathbf{x}} \in \mathbb{R}^n} f(\bar{\mathbf{x}}), \quad \text{subject to} \quad x_i \geq 0, \ i = 1, 2, \ldots, n$$

the trust-region subproblem may be

$$\min_{\bar{\mathbf{p}} \in \mathbb{R}^n} m_k(\bar{\mathbf{p}}), \quad \text{subject to} \quad \bar{\mathbf{x}}_k + \bar{\mathbf{p}} \geq 0, \text{ (component-wise)}, \|\bar{\mathbf{p}}\| \leq \Delta_k$$

This trust region is the intersection of the disk centered at $\bar{\mathbf{x}}_k$ and the first quadrant. It could look like this:

Such a region is hard to describe, and hard to work with.

If, instead, we work with the $\| \cdot \|_\infty$-norm, the trust region is the intersection of the square with sides Δ_k centered at $\bar{\mathbf{x}}_k$ and the first quadrant:

Much easier to work with...

Index

definition
    lim sup and lim inf, 15
lemma
    Cauchy point descent, 5
theorem
    Convergence (when $\eta = 0$), 13
    Convergence (when $\eta > 0$), 14
    Global trust-region Newton convergence ($\eta > 0$), 16
    Second order necessary conditions, 4

Reference(s):

MS-1983 J.J. Moré and D.C. Sorensen, *Computing a Trust Region Step*, SIAM Journal on Scientific and Statistical
         Computing, 4 (1983), pp. 553−572.