

Numerical Optimization

Lecture Notes #11

Conjugate Gradient Methods — Linear CG, Part #2

Fall 2022

Outline

- 1 Recap
 - Linear Conjugate Direction Methods
- 2 The Linear Conjugate Gradient Method
 - $A\bar{\mathbf{p}}$ vs. A^{-1}
 - Adding the Gradient to the Mix...
 - The CG Algorithm... and Krylov Subspaces
- 3 The CG Algorithm
 - Increased Efficiency
 - The Standard CG Algorithm
 - More Speed... Preconditioning
- 4 Non-Linear CG for Optimization Problems

Quick Recap: Linear Conjugate **Direction** Methods

We introduced the **conjugate direction method**: Given a starting point $\bar{\mathbf{x}}_0 \in \mathbb{R}^n$ and a set of conjugate directions $\{\bar{\mathbf{p}}_0, \bar{\mathbf{p}}_1, \dots, \bar{\mathbf{p}}_n\}$ the sequence

$$\bar{\mathbf{x}}_{k+1} = \bar{\mathbf{x}}_k + \alpha_k \bar{\mathbf{p}}_k, \quad \alpha_k = -\frac{\bar{\mathbf{r}}_k^T \bar{\mathbf{p}}_k}{\bar{\mathbf{p}}_k^T A \bar{\mathbf{p}}_k}, \quad \text{where } \bar{\mathbf{r}}_k = \bar{\mathbf{r}}(\bar{\mathbf{x}}_k) = A\bar{\mathbf{x}}_k - \bar{\mathbf{b}}$$

We showed that

- (1) this method is guaranteed to converge to the solution $\bar{\mathbf{x}}^* = A^{-1}\bar{\mathbf{b}}$ in at most n iterations;
- (2) each $\bar{\mathbf{x}}_k$ is the minimizer of $\Phi(\bar{\mathbf{x}}) = \frac{1}{2}\bar{\mathbf{x}}^T A \bar{\mathbf{x}} - \bar{\mathbf{b}}^T \bar{\mathbf{x}}$ over the set $\{\bar{\mathbf{x}} : \bar{\mathbf{x}} = \bar{\mathbf{x}}_0 + \text{span}\{\bar{\mathbf{p}}_0, \bar{\mathbf{p}}_1, \dots, \bar{\mathbf{p}}_{k-1}\}\}$

We're currently at the “**So what?!?**” stage...

More Questions than Answers

How can we make this useful?

- Given A , how do we get a set of conjugate vectors?
 - Efficiently, please!
 - Again, if we have the conjugate vectors, it seems like we will make more progress in certain directions than in others; hence, if we are planning on stopping short of n iterations, the subset of conjugate directions that we select will have an impact on how well we do...
 - Where is the gradient?

Comment: $A\bar{p}$ vs. A^{-1}





Figure: If the matrix A is sparse (many elements are zero, e.g. the matrix illustrated to the left), the computation of the matrix-vector product $A\bar{p}$ can be economized. However, generally, the inverse of a sparse matrix is dense (the matrix on the right). This is one indication that the conjugate direction method may be useful.

Enter: The Gradient

The **Conjugate Gradient** method is a conjugate direction method, which

- Generates the next conjugate vector \bar{p}_k using only the previous vector \bar{p}_{k-1} (earlier vectors are not needed.)
 - Cheap to compute, and store.
- Each direction \bar{p}_k is a clever linear combination of \bar{p}_{k-1} and the negative gradient of the objective $-\nabla\Phi(\bar{x}_k) = -\bar{r}(\bar{x}_k)$ (a.k.a “the (negative) residual,” or “the steepest descent direction.”)
 - Recall that we have a cheap update for the residual

$$\bar{r}_k = \bar{r}_{k-1} + \alpha_{k-1} \underbrace{A\bar{p}_{k-1}}_{\text{“Free”}}.$$

A New Conjugate Direction

We let the new conjugate direction be

$$\bar{\mathbf{p}}_k = -\bar{\mathbf{r}}_k + \beta_k \bar{\mathbf{p}}_{k-1},$$

and we select the scalar β_k so that $\bar{\mathbf{p}}_k$ and $\bar{\mathbf{p}}_{k-1}$ are A -conjugate

$$\bar{\mathbf{p}}_{k-1}^T A \bar{\mathbf{p}}_k = -\bar{\mathbf{p}}_{k-1}^T A \bar{\mathbf{r}}_k + \beta_k \bar{\mathbf{p}}_{k-1}^T A \bar{\mathbf{p}}_{k-1} = 0.$$

Hence,

$$\beta_k = \frac{\bar{\mathbf{p}}_{k-1}^T A \bar{\mathbf{r}}_k}{\bar{\mathbf{p}}_{k-1}^T A \bar{\mathbf{p}}_{k-1}} = \frac{\bar{\mathbf{r}}_k^T A \bar{\mathbf{p}}_{k-1}}{\bar{\mathbf{p}}_{k-1}^T A \bar{\mathbf{p}}_{k-1}},$$

where, again, the quantities $[A\bar{\mathbf{p}}_{k-1}]$ and $\bar{\mathbf{r}}_k^T A \bar{\mathbf{p}}_{k-1}$ are “free” (already computed).

Note!!! The first direction $\bar{\mathbf{p}}_0$ is set to be the steepest descent direction at the initial point $\bar{\mathbf{x}}_0$.

The Conjugate Gradient Algorithm (version 0.99 α)

Algorithm: Preliminary Conjugate Gradient

Given A , $\bar{\mathbf{b}}$ and $\bar{\mathbf{x}}_0$:

$\bar{\mathbf{r}}_0 = A\bar{\mathbf{x}}_0 - \bar{\mathbf{b}}$, $\bar{\mathbf{p}}_0 = -\bar{\mathbf{r}}_0$, $k = 0$

while ($\|\bar{\mathbf{r}}_k\| > 0$, or other stopping condition)

$$\alpha_k = -\frac{\bar{\mathbf{r}}_k^T \bar{\mathbf{p}}_k}{\bar{\mathbf{p}}_k^T A \bar{\mathbf{p}}_k},$$

Store the vector $A\bar{\mathbf{p}}_k$
and the scalar $\bar{\mathbf{p}}_k^T A \bar{\mathbf{p}}_k$

$$\bar{\mathbf{x}}_{k+1} = \bar{\mathbf{x}}_k + \alpha_k \bar{\mathbf{p}}_k$$

$$\bar{\mathbf{r}}_{k+1} = \bar{\mathbf{r}}_k + \alpha_k A \bar{\mathbf{p}}_k$$

$$\beta_{k+1} = \frac{\bar{\mathbf{r}}_{k+1}^T A \bar{\mathbf{p}}_k}{\bar{\mathbf{p}}_k^T A \bar{\mathbf{p}}_k}$$

$$\bar{\mathbf{p}}_{k+1} = -\bar{\mathbf{r}}_{k+1} + \beta_{k+1} \bar{\mathbf{p}}_k$$

$$k = k + 1$$

end-while

Does the CG Algorithm Work?

In order to guarantee convergence in n steps, the directions $\{\bar{p}_i\}$ must be A -conjugate; maybe we should show this?! But, first a definition:

Definition (Krylov Subspace)

A **Krylov** subspace of degree k for \bar{r}_0 is the space

$$\mathcal{K}(\bar{r}_0, k) \stackrel{\text{def}}{=} \text{span}\{\bar{r}_0, A\bar{r}_0, A^2\bar{r}_0, \dots, A^{k-1}\bar{r}_0\}.$$

We state a theorem which shows that the directions are indeed conjugate; further it shows that the residuals are mutually orthogonal, and that the search directions and residuals are contained in a Krylov subspace. These facts will allow us to optimize the CG algorithm for speed (computational effort).

“The CG Theorem”

Theorem

Suppose that the k^{th} iterate generated by the CG method is not the solution (i.e. $\bar{\mathbf{x}}_k \neq \bar{\mathbf{x}}^*$). The following properties hold

- (1) $\bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_i = 0, \quad i = 0, 1, \dots, k-1$
- (2) $\text{span}\{\bar{\mathbf{r}}_0, \bar{\mathbf{r}}_1, \dots, \bar{\mathbf{r}}_k\} = \text{span}\{\bar{\mathbf{r}}_0, A\bar{\mathbf{r}}_0, A^2\bar{\mathbf{r}}_0, \dots, A^k\bar{\mathbf{r}}_0\}$
- (3) $\text{span}\{\bar{\mathbf{p}}_0, \bar{\mathbf{p}}_1, \dots, \bar{\mathbf{p}}_k\} = \text{span}\{\bar{\mathbf{r}}_0, A\bar{\mathbf{r}}_0, A^2\bar{\mathbf{r}}_0, \dots, A^k\bar{\mathbf{r}}_0\}$
- (4) $\bar{\mathbf{p}}_k^T A\bar{\mathbf{p}}_i = 0, \quad i = 0, 1, \dots, k-1$

Therefore, the sequence $\{\bar{\mathbf{x}}_k\}$ converges to $\bar{\mathbf{x}}^*$ in at most n steps.

Note: The theorem is true **if and only if** the first search direction is the steepest descent direction. We notice that the search direction (and not the gradients/residuals) are conjugate in the “conjugate gradient method.”

A More Efficient Implementation

α_k (Numerator)

We now combine our results in order to tighten up the algorithm.

First, we use the relation (update for $\bar{\mathbf{p}}_k$ in the algorithm)

$$\bar{\mathbf{p}}_k = -\bar{\mathbf{r}}_k + \beta_k \bar{\mathbf{p}}_{k-1},$$

and the result (from lecture #10, or slide #10)

$$\bar{\mathbf{r}}_k^T \bar{\mathbf{p}}_i = 0, \quad i = 0, 1, \dots, k-1,$$

thus the numerator in the expression for α_k can be rewritten:

$$\bar{\mathbf{r}}_k^T \bar{\mathbf{p}}_k = \bar{\mathbf{r}}_k^T (-\bar{\mathbf{r}}_k + \beta_k \bar{\mathbf{p}}_{k-1}) = -\bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_k + \beta_k \underbrace{\bar{\mathbf{r}}_k^T \bar{\mathbf{p}}_{k-1}}_0 = -\bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_k.$$

A More Efficient Implementation

β_k

Second, we use the update formula for the residual

$$\bar{\mathbf{r}}_k = \bar{\mathbf{r}}_{k-1} + \alpha_{k-1} A \bar{\mathbf{p}}_{k-1} \quad \Leftrightarrow \quad \alpha_{k-1} A \bar{\mathbf{p}}_{k-1} = \bar{\mathbf{r}}_k - \bar{\mathbf{r}}_{k-1},$$

and again (from lecture #10, or slide #10)

$$\bar{\mathbf{r}}_k^T \bar{\mathbf{p}}_i = 0, \quad i = 0, 1, \dots, k-1,$$

as well as the update for $\bar{\mathbf{p}}_k$ in the algorithm

$$\bar{\mathbf{p}}_k = -\bar{\mathbf{r}}_k + \beta_k \bar{\mathbf{p}}_{k-1}$$

We get

$$\begin{aligned} \beta_k &= \frac{\bar{\mathbf{r}}_{k+1}^T A \bar{\mathbf{p}}_k}{\bar{\mathbf{p}}_k^T A \bar{\mathbf{p}}_k} = \frac{\bar{\mathbf{r}}_{k+1}^T (\bar{\mathbf{r}}_{k+1} - \bar{\mathbf{r}}_k)}{\alpha_k \bar{\mathbf{p}}_k^T A \bar{\mathbf{p}}_k} = \frac{\bar{\mathbf{r}}_{k+1}^T \bar{\mathbf{r}}_{k+1}}{(-\bar{\mathbf{r}}_k + \beta_k \bar{\mathbf{p}}_{k-1})^T (\bar{\mathbf{r}}_{k+1} - \bar{\mathbf{r}}_k)} \\ &= \frac{\bar{\mathbf{r}}_{k+1}^T \bar{\mathbf{r}}_{k+1}}{-\bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_{k+1} + \beta_k \bar{\mathbf{p}}_{k-1}^T \bar{\mathbf{r}}_{k+1} + \bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_k - \bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_{k+1}} = \frac{\bar{\mathbf{r}}_{k+1}^T \bar{\mathbf{r}}_{k+1}}{\bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_k}. \end{aligned}$$

The CG Algorithm (version 1.0, "Standard")

Algorithm: Conjugate Gradient

Given A , $\bar{\mathbf{b}}$ and $\bar{\mathbf{x}}_0$:

$$\bar{\mathbf{r}}_0 = A\bar{\mathbf{x}}_0 - \bar{\mathbf{b}}, \quad \bar{\mathbf{p}}_0 = -\bar{\mathbf{r}}_0, \quad k = 0$$

while ($\|\bar{\mathbf{r}}_k\| > 0$, or other stopping condition)

$$\alpha_k = \frac{\bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_k}{\bar{\mathbf{p}}_k^T A \bar{\mathbf{p}}_k},$$

Store the vector $A\bar{\mathbf{p}}_k$
 and the scalar $\bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_k$

$$\bar{\mathbf{x}}_{k+1} = \bar{\mathbf{x}}_k + \alpha_k \bar{\mathbf{p}}_k$$

$$\bar{\mathbf{r}}_{k+1} = \bar{\mathbf{r}}_k + \alpha_k A \bar{\mathbf{p}}_k$$

$$\beta_{k+1} = \frac{\bar{\mathbf{r}}_{k+1}^T \bar{\mathbf{r}}_{k+1}}{\bar{\mathbf{r}}_k^T \bar{\mathbf{r}}_k},$$

Keep numerator for next step!

$$\bar{\mathbf{p}}_{k+1} = -\bar{\mathbf{r}}_{k+1} + \beta_{k+1} \bar{\mathbf{p}}_k$$

$$k = k + 1$$

end-while

The CG Algorithm (version 1.0, “Standard”)

Work

The work per iteration for this version of the CG algorithm consists of

- One matrix-vector product $A\bar{\mathbf{p}}_k$
 $\sim n^2$ operations (if A is dense)
 $\sim \mathcal{O}(n)$ in many cases, when A is sparse.
- Two inner products: $\bar{\mathbf{p}}_k^T (A\bar{\mathbf{p}}_k)$ and $\bar{\mathbf{r}}_{k+1}^T \bar{\mathbf{r}}_{k+1}$
 $\sim n$ additions, and $\sim n$ multiplications
- Three vector sums
 $\sim 3n$ additions

The CG Algorithm — Convergence

In **exact arithmetic** CG converges in at most n iterations.

In many cases, the algorithm will find the solution in many fewer iterations. We leave the detailed convergence analysis for some other day, but state some key results:

Theorem

If A has only r distinct eigenvalues, then the CG iteration will terminate at the solution $\bar{\mathbf{x}}^$ in at most r iterations.*

Theorem

If A has eigenvalues $\lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$, we have that

$$\|\bar{\mathbf{x}}_{k+1} - \bar{\mathbf{x}}^*\|_A \leq \left[\frac{\lambda_{n-k} - \lambda_1}{\lambda_{n-k} + \lambda_1} \right] \|\bar{\mathbf{x}}_0 - \bar{\mathbf{x}}^*\|_A.$$

The CG Algorithm — Convergence: Comments

The second theorem tells us that the CG algorithm selects **exactly** the optimal sequence of conjugate search directions $\{\bar{\mathbf{p}}_i\}$.

If there is a cluster of eigenvalues of A around λ_1 , i.e. $\lambda_1 = 1$, $\lambda_{3900} = 1.0002$, $\lambda_{4000} = 1.03$, and $\lambda_n = \lambda_{4032}$, then after 32 iterations we would have

$$\|\bar{\mathbf{x}}_{32} - \bar{\mathbf{x}}^*\|_A \leq \left[\frac{0.03}{2.03} \right] \|\bar{\mathbf{x}}_0 - \bar{\mathbf{x}}^*\|_A$$

and after another 100 iterations

$$\|\bar{\mathbf{x}}_{132} - \bar{\mathbf{x}}^*\|_A \leq \left[\frac{0.0002}{2.0002} \right] \|\bar{\mathbf{x}}_0 - \bar{\mathbf{x}}^*\|_A$$

With tight clustering (which is quite common) we often achieve very good convergence after a $k \ll n$ iterations.

The CG Algorithm — Example

2 of 2

Let

$$A = \text{diag}(k^2 I_k, k = 1 \dots 5), \quad \bar{\mathbf{b}} = \bar{\mathbf{1}}, \quad \bar{\mathbf{x}}_0 = \bar{\mathbf{0}}.$$

We get $\bar{\mathbf{x}}_1, \bar{\mathbf{x}}_2, \bar{\mathbf{x}}_3, \bar{\mathbf{x}}_4, \bar{\mathbf{x}}_5 = \bar{\mathbf{x}}^*$:

$$\begin{bmatrix} 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \\ 0.0667 \end{bmatrix}, \begin{bmatrix} 0.2 \\ 0.179 \\ 0.179 \\ 0.143 \\ 0.143 \\ 0.143 \\ 0.0929 \\ 0.0929 \\ 0.0929 \\ 0.0929 \\ 0.0286 \\ 0.0286 \\ 0.0286 \\ 0.0286 \\ 0.0286 \end{bmatrix}, \begin{bmatrix} 0.4 \\ 0.293 \\ 0.293 \\ 0.154 \\ 0.154 \\ 0.154 \\ 0.0429 \\ 0.0429 \\ 0.0429 \\ 0.0429 \\ 0.0429 \\ 0.0429 \\ 0.0429 \\ 0.0429 \\ 0.0429 \end{bmatrix}, \begin{bmatrix} 0.667 \\ 0.345 \\ 0.345 \\ 0.0873 \\ 0.0873 \\ 0.0873 \\ 0.0665 \\ 0.0665 \\ 0.0665 \\ 0.0665 \\ 0.0397 \\ 0.0397 \\ 0.0397 \\ 0.0397 \\ 0.0397 \end{bmatrix}, \begin{bmatrix} 1 \\ 0.25 \\ 0.25 \\ 0.111 \\ 0.111 \\ 0.111 \\ 0.0625 \\ 0.0625 \\ 0.0625 \\ 0.0625 \\ 0.04 \\ 0.04 \\ 0.04 \\ 0.04 \\ 0.04 \end{bmatrix}$$

Contrasting Example: “Maximal Subspace Collapse”

1 of 2

What happens is we select a sequence of search directions which collapse the maximal number of dimensions of the residual (which coincidentally, in this example, is the subspace corresponding to the largest eigenvalue)?

$$\bar{\mathbf{p}}_0 = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix}, \quad \bar{\mathbf{p}}_1 = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \bar{\mathbf{p}}_2 = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \bar{\mathbf{p}}_3 = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \bar{\mathbf{p}}_4 = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Contrasting Example: “Maximal Subspace Collapse”

2 of 2

We get the following sequence of residuals

$$\bar{\mathbf{r}}_1^{\text{MSC}} \approx \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \bar{\mathbf{r}}_2^{\text{MSC}} \approx \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \bar{\mathbf{r}}_3^{\text{MSC}} \approx \begin{bmatrix} 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \bar{\mathbf{r}}_4^{\text{MSC}} \approx \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \bar{\mathbf{r}}_5^{\text{MSC}} \approx \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

	$\ \bar{\mathbf{r}}_0\ $	$\ \bar{\mathbf{r}}_1\ $	$\ \bar{\mathbf{r}}_2\ $	$\ \bar{\mathbf{r}}_3\ $	$\ \bar{\mathbf{r}}_4\ $	$\ \bar{\mathbf{r}}_5\ $
CG	$\sqrt{15}$	2.1603	1.5492	1.1339	0.7454	≈ 0
MSC	$\sqrt{15}$	3.1623	2.4495	1.7321	1.0000	≈ 0

Table: CG gives the optimal sequence of residual lengths at each iteration!

Contrasting Example #2: Steepest Descent




Figure: Since the condition number is only $5^2 = 25$, one may think that maybe steepest descent will do a decent job? But, alas, it takes over 200 iterations to get a reduction of the residual norm by 10^{-8} .

The CG Algorithm — Convergence: More Comments

It is worth noting that the theorem gives an **upper bound** of the error, in practice it is *almost* true that if the eigenvalues of A occur in r distinct clusters, then (compare with the first theorem) the CG algorithm will approximately solve the problem after r steps.

Further it can be shown that for the CG algorithm

$$\|\bar{\mathbf{x}}_k - \bar{\mathbf{x}}^*\|_A \leq 2 \left[\frac{\sqrt{\kappa(A)} - 1}{\sqrt{\kappa(A)} + 1} \right]^k \|\bar{\mathbf{x}}_0 - \bar{\mathbf{x}}^*\|_A,$$

whereas (forgotten from lecture #5)

$$\|\bar{\mathbf{x}}_{k+1} - \bar{\mathbf{x}}^*\|_A \leq \left[\frac{\kappa(A) - 1}{\kappa(A) + 1} \right]^k \|\bar{\mathbf{x}}_0 - \bar{\mathbf{x}}^*\|_A,$$

for the steepest descent algorithm. Here $\kappa(A) = \lambda_n/\lambda_1$ is the condition number of the matrix A .

CG vs. Steepest Descent

1 of 2




Figure: Comparing the convergence factors $\left[\frac{\sqrt{\kappa(A)}-1}{\sqrt{\kappa(A)}+1} \right]$ (for Conjugate Gradient) and $\left[\frac{\kappa(A)-1}{\kappa(A)+1} \right]$ (for Steepest Descent) for condition numbers, $\kappa(A) \in [2, 200]$.

CG vs. Steepest Descent

2 of 2

$\kappa(A)$	CG	SD	SD/CG
10	22	69	3.136
100	72	691	9.597
1,000	229	6,908	30.166
10,000	725	69,078	95.280

Table: A comparison of how many iterations Conjugate gradient (CG) and Steepest descent (SD) are required in order to reduce the initial error $\|\bar{\mathbf{x}}_0 - \bar{\mathbf{x}}^*\|_A$ by a factor of 10^{-6} . We notice a “slight” improvement. The speedup is $\sim \sqrt{\kappa(A)}$.

Preconditioning: Speeding Things Up Even More

The CG method can be accelerated further by **preconditioning** the linear system; we [FORMALLY] make a non-singular change of variables

$$\hat{\mathbf{x}} = C\bar{\mathbf{x}},$$

and solve the linear system (and/or its equivalent minimization problem, $\min \hat{\Phi}(\bar{\mathbf{x}})$)

$$[C^{-T}AC^{-1}]\hat{\mathbf{x}} - C^{-T}\bar{\mathbf{b}}, \quad \hat{\Phi}(\bar{\mathbf{x}}) = \frac{1}{2}\hat{\mathbf{x}}^T [C^{-T}AC^{-1}]\hat{\mathbf{x}} - [C^{-T}\bar{\mathbf{b}}]^T \hat{\mathbf{x}}.$$

Now, the convergence rate will depend on the eigenvalues of $\mathcal{A} = [C^{-T}AC^{-1}]$. Therefore, we would like to choose C such that the eigenvalues of \mathcal{A} are favorably clustered, and/or the condition number of \mathcal{A} is less than that of A .

As in the case of the transformation guaranteeing n -step convergence, this change of variables does not have to be done explicitly.

The Preconditioned CG Algorithm (a.k.a. "PCG")

Algorithm: PCG

Given A , $\mathbf{M} = \mathbf{C}^T \mathbf{C}$, $\bar{\mathbf{b}}$ and $\bar{\mathbf{x}}_0$: compute $\bar{\mathbf{r}}_0 = A\bar{\mathbf{x}}_0 - \bar{\mathbf{b}}$,
 $\bar{\mathbf{y}}_0 = \mathbf{M}^{-1}\bar{\mathbf{r}}_0$, $\bar{\mathbf{p}}_0 = -\bar{\mathbf{y}}_0$, $k = 0$

while ($\|\mathbf{r}_k\| > 0$, or other stopping condition)

$$\alpha_k = \frac{\bar{\mathbf{r}}_k^T \bar{\mathbf{y}}_k}{\bar{\mathbf{p}}_k^T A \bar{\mathbf{p}}_k},$$

Store the vector $A\bar{\mathbf{p}}_k$
 and the scalar $\bar{\mathbf{r}}_k^T \bar{\mathbf{y}}_k$

$$\bar{\mathbf{x}}_{k+1} = \bar{\mathbf{x}}_k + \alpha_k \bar{\mathbf{p}}_k$$

$$\bar{\mathbf{r}}_{k+1} = \bar{\mathbf{r}}_k + \alpha_k A \bar{\mathbf{p}}_k$$

$$\bar{\mathbf{y}}_{k+1} = \mathbf{M}^{-1} \bar{\mathbf{r}}_{k+1}$$

$$\beta_{k+1} = \frac{\bar{\mathbf{r}}_{k+1}^T \bar{\mathbf{y}}_{k+1}}{\bar{\mathbf{r}}_k^T \bar{\mathbf{y}}_k},$$

Save the numerator for next step!

$$\bar{\mathbf{p}}_{k+1} = -\bar{\mathbf{y}}_{k+1} + \beta_{k+1} \bar{\mathbf{p}}_k$$

$$k = k + 1$$

end-while

PCG: Comments

If we set $M = I$, then we recover standard CG.

We note that in each iteration, we have to solve the linear system $M\bar{\mathbf{y}}_{k+1} = \bar{\mathbf{r}}_{k+1} \Leftrightarrow (\bar{\mathbf{y}}_{k+1} = M^{-1}\bar{\mathbf{r}}_{k+1})$. We must select M so that we can do this quickly, otherwise we lose the overall-work advantage over Steepest descent or Gaussian elimination.

There are several (usually competing) properties we would like M to have:

- M should effectively impact the structure of the eigenvalues.
- M should be cheap to compute and store.
- The linear system $M\bar{\mathbf{y}} = \bar{\mathbf{r}}$ should be “easy.”

Finding Good Preconditions

1 of 2

There is no “best” way of finding M . The optimal M for a particular A may even depend on how much memory, etc your computer has.

In general M is a simplified version of A , e.g. we may take the tridiagonal part of A :




Figure: When A has a banded structure (left) with a significant bandwidth, then a tri-diagonal preconditioner M (right) may be a good choice. Recall that in this case we can solve $M\bar{y} = \bar{r}$ in $\mathcal{O}(n)$ operations.

Finding Good Preconditioners

2 of 2

Preconditioning is itself a science (or an art?) which will be revisited in more detail in (???)

One of the more efficient strategies is **incomplete Cholesky factorization**. — The exact Cholesky factorization of an SPD matrix A has the form

$$LL^T = A, \quad \text{where } L \text{ is lower triangular.}$$

Usually, even though A may be sparse, L will be dense (due to **fill-ins**). In incomplete Cholesky factorization, the same algorithm is followed, but whenever a fill-in occurs, that value is dropped — this way we end up with

$$M = \tilde{L}\tilde{L}^T \approx A,$$

where \tilde{L} and A have the same sparsity patterns.

Index

algorithm

conjugate gradient, 13

preconditioned conjugate gradient, 26

Krylov subspace, 9