Numerical Optimization

Lecture Notes #13 — Practical Newton Methods
Inexact Newton Steps / Linesearch Newton Methods

Fall 2022

Quick Recap: Nonlinear Conjugate Gradient Methods

Extension of the linear CG to work for non-linear (optimization) problems.

In the first pass (Fletcher-Reeves' Algorithm), we simply replaced all instances of the residual $\bar{\mathbf{r}}_k$ by the gradient of the objective $\nabla f(\bar{\mathbf{x}}_k)$, and the step length α_k is calculated by a linesearch.

We looked at some modifications, and arrived at the Polak-Ribière PR+CG algorithm, where the β of Fletcher-Reeves is modified

$$\beta_{k+1}^{\mathrm{FR}} = \frac{\nabla f_{k+1}^T \nabla f_{k+1}}{\nabla f_k^T \nabla f_k} \quad \rightarrow \quad \beta_{\mathbf{k+1}}^{\mathrm{PR}} = \frac{\nabla \mathbf{f_{k+1}^T}(\nabla \mathbf{f_{k+1}} - \nabla \mathbf{f_k})}{\nabla \mathbf{f_k^T} \nabla \mathbf{f_k}}$$

and the final β is $\beta_{k+1}^+ = \max(\beta_{k+1}^{\mathrm{PR}}, 0)$.

Finally, periodic restarting, when

$$\frac{|\nabla f_k^T \nabla f_{k-1}|}{\nabla f_k^T \nabla f_k} \geq \nu \sim 0.1$$

was introduced in order to ensure good convergence.

Practical Newton Methods: Introduction

We know (*e.g.* from LECTURE #5) that Newton's method has great local convergence properties. Once we get close to the minimizer $\bar{\mathbf{x}}^*$ convergence is **quadratic**.

This convergence requires that we start "close enough" to $\bar{\mathbf{x}}^*$ — in regions far away, where the objective is not convex, all bets are off and the behavior can be quite erratic; we cannot guarantee convergence at all!

Our present goal:
⇒ To design a Newton-based method which is robust and efficient in "all" cases.

The Newton Step

We get the Newton step from the symmetric $n \times n$ linear system

The Newton Direction: $\qquad \nabla^2 f(\bar{\mathbf{x}}_k)\bar{\mathbf{p}}_k^N = -\nabla f(\bar{\mathbf{x}}_k)$

For **global convergence** the Newton direction must be a **descent direction**, this is true if the Hessian $(\nabla^2 f(\bar{\mathbf{x}}_k))$ is **positive definite**.

If the Hessian is not positive definite, the Newton direction may be an **ascent direction** and/or extremely long (division by almost zero).

We look at two approaches: The first uses the conjugate gradient method, and gives us the **"Newton-CG"** methods for both line-search and trust-region methods; the second strategy involves modifying the Hessian so that it becomes *"sufficiently positive definite,"* yielding the **"modified Newton method."**

Computational Cost

As always, we want to keep the computational cost down.

In Newton-CG, this is accomplished by terminating the computation before an exact solution to

$$\nabla^2 f(\bar{\mathbf{x}}_k)\bar{\mathbf{p}}_k^N = -\nabla f(\bar{\mathbf{x}}_k),$$

has been found. Thus we get and **approximation** $\bar{\mathbf{p}}_k \approx \bar{\mathbf{p}}_k^N$, hence the name **"inexact Newton methods."**

We would like to exploit any special **sparsity structure** in the Hessian in order to solve the linear problem as efficiently as possible.

For now, we assume we **have** access to the Hessian in **analytical** form. We will cover this final issue soon.

Inexact Newton Steps

If we settle for an inexact (approximate) solution of

$$\nabla^2 f(\bar{\mathbf{x}}_k)\bar{\mathbf{p}}_k^N = -\nabla f(\bar{\mathbf{x}}_k),$$

then we need a measure of how close our approximate solution $\bar{\mathbf{p}}_k$ is to the exact Newton direction... Another use of the **residual**

$$\bar{\mathbf{r}}_k = \nabla^2 f(\bar{\mathbf{x}}_k)\bar{\mathbf{p}}_k + \nabla f(\bar{\mathbf{x}}_k).$$

Usually we do not want the termination condition for the inexact solution to depend on the size of f, hence are interested in the **relative size** of the residual, and say that an approximate solution is good enough when

Termination criterion: $\qquad \|\bar{\mathbf{r}}_k\| \leq \eta_k \|\nabla f(\bar{\mathbf{x}}_k)\|, \quad \eta_k \in (0, 1)$

The sequence $\{\eta_k\}$ is known as a **forcing sequence**.

The Forcing Sequence: Convergence

Selection of the forcing sequence greatly impacts how fast the overall algorithm converges, as illustrated in the following results:

Theorem

Suppose that the gradient $\nabla f(\bar{\mathbf{x}})$ is continuously differentiable in a neighborhood $\mathcal{N}$ of a minimizer $\bar{\mathbf{x}}^*$, and assume that the Hessian $\nabla^2 f(\bar{\mathbf{x}}^*)$ is positive definite. Consider the iteration $\bar{\mathbf{x}}_{k+1} = \bar{\mathbf{x}}_k + \bar{\mathbf{p}}_k$ where $\bar{\mathbf{r}}_k(\bar{\mathbf{p}}_k)$ satisfies

$$\|\bar{\mathbf{r}}_k\| \leq \eta_k \|\nabla f(\bar{\mathbf{x}}_k)\|, \quad \eta_k \in (0, 1)$$

and assume that $\eta_k \leq \eta$ for some $\eta \in [0, 1)$. Then, if the starting point $\bar{\mathbf{x}}_0$ is sufficiently near $\bar{\mathbf{x}}^*$, the sequence $\{\bar{\mathbf{x}}_k\}$ converges to $\bar{\mathbf{x}}^*$ linearly. That is, for all sufficiently large k we have

$$\|\bar{\mathbf{x}}_{k+1} - \bar{\mathbf{x}}^*\| \leq c \|\bar{\mathbf{x}}_k - \bar{\mathbf{x}}^*\|$$

for some constant $c \in (0, 1)$.

Practical Newton Methods
Line Search Newton-CG Method
Modified Newton's Method

Introduction
Newton-CG — "Inexact Newton Methods"

The Forcing Sequence: Convergence   2 of 3

The preceding theorem is neither very exciting, nor very useful (on its own).

The restriction on the forcing sequence is very mild, we are basically just requiring that we make *some* progress in solving the linear system

$$\nabla^2 f(\bar{\mathbf{x}}_k)\bar{\mathbf{p}}_k^N = -\nabla f(\bar{\mathbf{x}}_k).$$

Likewise, the result — **linear convergence** — is good news, but hardly anything that causes us to throw a party!

However, by **carefully selecting** the forcing sequence we get a slightly more exciting result...

Practical Newton Methods
Line Search Newton-CG Method
Modified Newton's Method

Introduction
Newton-CG — "Inexact Newton Methods"

The Forcing Sequence: Convergence   3 of 3

Theorem

*Suppose that the conditions of the previous theorem hold, and
assume that the iterates $\{\bar{\mathbf{x}}_k\}$ generated by the inexact Newton
method converge to $\bar{\mathbf{x}}^*$. Then the rate of convergence is*
superlinear *if $\eta_k \to 0$, and* **quadratic** *if $\eta_k = \mathcal{O}(\|\nabla f(\bar{\mathbf{x}}_k)\|)$.*

Now we know exactly how hard we have to work at solving the
linear systems in order to achieve certain convergence rates, *e.g.*

$$\eta_k = \min\left(10^{-3}, \sqrt{\|\nabla f(\bar{\mathbf{x}}_k)\|}\right) \quad \textbf{Superlinear Convergence}$$

$$\eta_k = \min\left(10^{-3}, \|\nabla f(\bar{\mathbf{x}}_k)\|\right) \quad \textbf{Quadratic Convergence}$$

Note: These results are still **local** — we still have to figure out how
to make our algorithms work if not started "close" to $\bar{\mathbf{x}}^*$.

We now have the pieces necessary to build **robust Newton methods** with good performance characteristics: first on the menu — the line search Newton-CG method:

Getting the search direction $\bar{\mathbf{p}}_k$:

We apply the linear Conjugate Gradient (CG) method to the Newton equations

$$\nabla^2 f(\bar{\mathbf{x}}_k)\bar{\mathbf{p}}_k^N = -\nabla f(\bar{\mathbf{x}}_k),$$

and require that the solution satisfies a termination test of the type

$$\|\bar{\mathbf{r}}_k\| \leq \eta_k \|\nabla f(\bar{\mathbf{x}}_k)\|, \quad \eta_k \in (0, 1).$$

However, if the Hessian is not positive definite this may break...

If/When the Hessian is not positive definite we may enter a region of negative curvature; when we do, the CG iteration is terminated in order to guarantee that the generated $\bar{\mathbf{p}}_k$ is a descent direction:

In line-direction-search, we set $A = \nabla^2 f(\bar{\mathbf{x}}_k)$, $\bar{\mathbf{b}} = -\nabla f(\bar{\mathbf{x}}_k)$ and then start the CG-iteration:

(1)   The starting point is set to $\bar{\mathbf{x}}^{(0)} = 0$

(2)   If a (**CG-internal**) search direction $\bar{\mathbf{p}}^{(i)}$ generated by the CG-iteration satisfies

$$\left[\bar{\mathbf{p}}^{(i)}\right]^T A \left[\bar{\mathbf{p}}^{(i)}\right] \leq 0, \quad \textbf{Negative curvature test}$$

then, if ($i == 0$), set $\bar{\mathbf{x}}^{(0)} = \bar{\mathbf{b}} = -\nabla f(\bar{\mathbf{x}}_k)$ [Steepest Descent] and return, otherwise stop immediately and return $\bar{\mathbf{x}}^{(i)}$.

(3)   The approximate Newton step $\bar{\mathbf{p}}_k \stackrel{\text{def}}{=} \bar{\mathbf{x}}^{(i)}$.

Line Search Newton-CG Method

Algorithm: "CG-core"

```
Given A, b̄, η^(κ), x̄₀^CG: r̄₀ = Ax̄₀^CG − b̄, p̄₀ = −r̄₀, k = 0
while ( ‖r̄ₖ‖ > η^(κ)‖b̄‖ )
```

$$\alpha_k = \frac{\bar{r}_k^T \bar{r}_k}{\bar{p}_k^T A \bar{p}_k},$$
Store the vector $A\bar{p}_k$ and the scalar $\bar{r}_k^T \bar{r}_k$

$$\text{if} \quad \bar{p}_k^T A \bar{p}_k \leq 0, \ k > 0 \quad \text{return}(\bar{x}_k^{CG})$$

$$\text{if} \quad \bar{p}_k^T A \bar{p}_k \leq 0, \ k = 0 \quad \text{return}(\bar{p}_0)$$

$$\bar{x}_{k+1}^{CG} = \bar{x}_k^{CG} + \alpha_k \bar{p}_k$$

$$\bar{r}_{k+1} = \bar{r}_k + \alpha_k A\bar{p}_k$$

$$\beta_{k+1} = \frac{\bar{r}_{k+1}^T \bar{r}_{k+1}}{\bar{r}_k^T \bar{r}_k},$$
Save numerator for next iteration!

$$\bar{p}_{k+1} = -\bar{r}_{k+1} + \beta_{k+1} \bar{p}_k$$

```
end-while( k = k + 1 )
```

—— The $\bar{p}_k$'s are CG-internal search directions, **not** to be confused with the search direction for the optimization algorithm!

Line Search Newton-CG Method

Note: If the CG-core algorithm encounters a **direction of negative curvature** in the first iteration, the steepest descent direction is used.

Algorithm: Line Search Newton-CG Method

```
Given x̄₀:   k = 0
while ( x̄_k is not a minimum, e.g.   ‖∇f(x̄_k)‖ ≥ 10⁻⁶ )
       p̄_k^N-CG   =   CG-core(A = ∇²f(x̄_k), b̄ = -∇f(x̄_k), η^(κ) = η_k, x̄₀^CG = 0̄)
       α_k^LS    =   linesearch_Strong_Wolfe(p̄_k^N-CG,...)
       x̄_{k+1}   =   x̄_k + α_k^LS p̄_k^N-CG
end-while ( k = k + 1 )
```

Where we specify η_k as discussed earlier, and the linesearch is such that α_k satisfies the Wolfe, Strong Wolfe, Goldstein, or Armijo backtracking conditions.

Comments:

- Nothing is stopping us from basing this on the preconditioned version of CG, in fact that is probably the right thing to do

- Line Search Newton-CG (LS-N-CG) is well suited for large problems.

- LS-N-CG has one minor weakness — If/When the Hessian is nearly singular, the Newton-CG direction can be excessively long resulting in many function evaluations in the linesearch.

- This weakness is greatly alleviated by preconditioning, *i.e.* implementing LS-N-PCG(M).

Sometimes it is desirable to use a direct linear algebra technique, *i.e.* an efficient cousin of Gaussian Elimination, to solve the Newton equations

$$\nabla^2 f(\bar{\mathbf{x}}_k)\bar{\mathbf{p}}_k^N = -\nabla f(\bar{\mathbf{x}}_k).$$

If/When the Hessian is not positive definite (or close to singular), it can be modified either before or **during** the solution process so that in effect we solve

$$\underbrace{\left[\nabla^2 f(\bar{\mathbf{x}}_k) + E_k\right]}_{\text{Sufficiently Positive Definite}} \bar{\mathbf{p}}_k^N = -\nabla f(\bar{\mathbf{x}}_k),$$

where the Hessian modification E_k is chosen so that the resulting matrix is sufficiently positive definite.

Algorithm: Line Search Newton with Modification

```
Given x̄₀:  k = 0
while ( x̄ₖ is not a minimum, e.g.  ‖∇f(x̄ₖ)‖ ≥ 10⁻⁶ )
```

$$B_k = \texttt{factorize}(\nabla^2 f(\bar{\mathbf{x}}_k) + E_k)$$

$$\bar{\mathbf{p}}_k^{\texttt{N-mod}} = -B_k^{-1}\nabla f(\bar{\mathbf{x}}_k)$$

$$\alpha_k^{\texttt{LS}} = \texttt{linesearch\_Strong\_Wolfe}(\bar{\mathbf{p}}_k^{\texttt{N-mod}}, \dots)$$

$$\bar{\mathbf{x}}_{k+1} = \bar{\mathbf{x}}_k + \alpha_k^{\texttt{LS}}\bar{\mathbf{p}}_k^{\texttt{N-mod}}$$

```
end-while ( k = k + 1 )
```

Where the linesearch is such that α_k satisfies the Wolfe, Strong Wolfe, Goldstein, or Armijo backtracking conditions.

The **factorization** algorithm is such that $E_k = 0$ if $\nabla^2 f(\bar{\mathbf{x}}_k)$ is sufficiently positive definite; otherwise chosen so that $B_k = (\nabla^2 f(\bar{\mathbf{x}}_k) + E_k)$ is sufficiently positive definite.

We save the details of Hessian modification for next lecture...

Index